

# Machine Learning (CE 40717)

## Fall 2024

Ali Sharifi-Zarchi

CE Department  
Sharif University of Technology

September 20, 2024



- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs
- 4 Applications
- 5 Shortcomings
- 6 Conclusion
- 7 References

# 1 Introduction

## 2 Principal Component Analysis (PCA)

## 3 Choose PCs

## 4 Applications

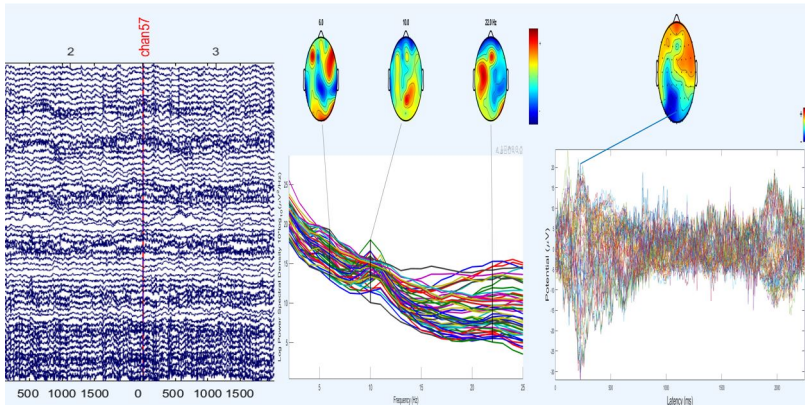
## 5 Shortcomings

## 6 Conclusion

## 7 References

# High Dimensional Data

- High-Dimensions = Lots of Features
- EEG Signals of Brain 64 Channels \* 3000 Time Points For Each Trial




# High Dimensional Data

- High-Dimensions = Lots of Features
- High Resolution Images (Millions of Pixels)




# Dimensionality Reduction

- **Feature Selection**
  - Select a subset of a given feature set
- **Feature Extraction**
  - A linear or non-linear transform on the original feature space



feature extraction



feature selection

# Dimensionality Reduction Benefits

- **Visualization**
  - Project high dimensional data into 2D or 3D
- **More efficient use of resources**
  - Time, Memory, CPU
- **Statistical**
  - Fewer dimensions leads to better generalization
- **Removing Noise**
- **Pre-Process**
  - Improve accuracy by reducing features
  - As a Preprocessing step to reduce dimensions for supervised learning tasks
  - Helps avoiding overfitting

## 1 Introduction

## 2 Principal Component Analysis (PCA)

## Sequential Algorithm

## Sample Covariance Matrix Algorithm

### 3 Choose PCs

## 4 Applications

## 5 Shortcomings

## 6 Conclusion

## 7 References





# Definition

- **Goal** is reducing the dimensionality of the data while preserving important aspects of the data
- **Principal Components (PCs)** are orthogonal vectors that are ordered by the fraction of the total information (variation) in the corresponding directions
  - Find the directions at which data approximately lie

# Definition


Orthogonal projection of the data onto a **lower-dimensional** linear space that:

- Maximizes variance of projected data
- Minimizes sum of squared distances to the line



# Definition

- Minimizing sum of square distances to the line is **equivalent** to maximizing the sum of squares of the projections on that line.




# Idea

- Given data points in a  $d$ -dimensional space, project them into a lower dimensional space while preserving as much information as possible,
  - Find best planar approximation of 3D data
  - Find best 12-D approximation of 104-D data
- In particular, choose projection that minimizes squared error in reconstructing the original data


# Idea

- 2D Gaussian dataset




# Idea

- 2D Gaussian dataset
- First PCA axis




# Idea

- 2D Gaussian dataset
- First and second PCA axes



# Random vs Principal Projection

- Random direction vs. principal component






# Pre-processing

- **Center the data**
  - **Zeroing** out the **mean** of each feature
- **Scaling to normalize each feature to have variance 1**
  - An arbitrary step (May affect the final result!)
  - It helps when unit of measurements of features are different and some features may be ignored without normalization

# Pre-processing

- Scaling to normalize each feature may affect the final result!!



# Algorithms

- Algorithm 1: sequential
- Algorithm 2: sample covariance matrix

## 1 Introduction

## 2 Principal Component Analysis (PCA)

## Sequential Algorithm

## Sample Covariance Matrix Algorithm

### 3 Choose PCs

## 4 Applications

## 5 Shortcomings


## 6 Conclusion

## 7 References

# Sequential Algorithm

- First view
  - Find directions with the maximum variations

$$\max_{v_1} \frac{1}{N} \sum_{n=1}^N (v_1^T x_n)^2 = \frac{1}{N} \sum_{n=1}^N v_1^T (x_n x_n^T) v_1 = v_1^T \left( \frac{1}{N} \sum_{n=1}^N (x_n x_n^T) \right) v_1 = v_1^T S v_1$$
$$\text{s.t. } v_1^T v_1 = 1$$



# Sequential Algorithm

- First view
- why  $v_1^T v_1 = 1$ ?
- Eigenvector with maximum eigenvalue maximizes the objective
  - Using Lagrangian multiplier technique

$$L(v_1, \lambda) = v_1^T S v_1 + \lambda(1 - v_1^T v_1)$$

$$\frac{\partial L}{\partial v_1} = 0 \Rightarrow 2Sv_1 - 2\lambda v_1 = 0$$

$$\Rightarrow Sv_1 = \lambda v_1$$

# Sequential Algorithm

- First view
- To find  $v_2$ , we maximize the variance of the projection in the residual subspace

$$v_2 = \max_{v_2} \left( \frac{1}{N} \sum_{i=1}^N (x_i - v_1^T x_i)^2 \right)$$

$$\text{s.t. } v_2^T v_2 = 1$$

- To find  $v_k$ , we maximize the variance of the projection in the residual subspace

$$v_k = \max_{v_k} \left( \frac{1}{N} \sum_{i=1}^N \left( x_i - \sum_{j=1}^{k-1} W_j^T x_i \right)^2 \right)$$


$$\text{s.t. } v_k^T v_k = 1$$

# Sequential Algorithm

- Second view
  - Find directions with the minimum reconstruction error

$$\min_{v_1} \sum_{n=1}^N \|x_n - (v_1^T x_n) v_1\|_2^2$$
$$\text{s.t. } v_1^T v_1 = 1$$

- Show this has an equal solution with the first view





# Sequential Algorithm

- As we have  $Sv_j = \lambda_j v_j$ ,

$$\Rightarrow \text{var}(v_j^T x) = v_j^T x x^T v_j = v_j^T S v_j = \lambda_j v_j^T v_j = \lambda_j$$

.

- The variance along an eigenvector  $v_j$  equals the eigenvalue  $\lambda_j$ .

# Sequential Algorithm

- Eigenvalues:  $\lambda_1 \geq \lambda_2 \geq \lambda_3 \geq \dots$ 
  - The first PC  $\nu_1$  is the the eigenvector of the sample covariance matrix  $S$  associated with the largest eigenvalue
  - The 2nd PC  $\nu_2$  is the the eigenvector of the sample covariance matrix  $S$  associated with the second largest eigenvalue
  - And so on ...
- To reduce the dimension of the data to  $k$ , we select eigenvectors with the top  $k$  eigenvalues

## 1 Introduction

## 2 Principal Component Analysis (PCA)

## Sequential Algorithm

## Sample Covariance Matrix Algorithm

### 3 Choose PCs

## ④ Applications

## 5 Shortcomings

## 6 Conclusion

## 7 References

# Sample Covariance Matrix

- Given data  $x_1, \dots, x_n$ , compute covariance matrix  $\Sigma$

$$\Sigma = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(x_i - \bar{x})^T \text{ where } \bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$$

- PCA basis vectors = the eigenvectors of  $\Sigma$
- Larger eigenvalue  $\rightarrow$  more important eigenvectors

# Sample Covariance Matrix

---

## Algorithm 1 Sample Covariance Matrix

---

- 1: **Input:**  $X \in \mathbb{R}^{N \times d}$  (data matrix with  $N$  data points and  $d$  dimensions)
  - 2: Compute the mean of each feature:  $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$
  - 3: Subtract the mean from each data point (center the data):  $\tilde{X} \leftarrow X - 1_N \bar{x}^T$
  - 4: Compute the covariance matrix:  $S = \frac{1}{N} \tilde{X}^T \tilde{X}$
  - 5: Compute the eigenvalues and eigenvectors of  $S$ :  $[\lambda_1, \lambda_2, \dots, \lambda_d], [v_1, v_2, \dots, v_d] = \text{eig}(S)$
  - 6: Select the top  $K$  eigenvectors corresponding to the largest eigenvalues:  $A \leftarrow [v_1, v_2, \dots, v_K]$
  - 7: Transform the data into the new subspace:  $X' \leftarrow X \cdot A$
  - 8: **Output:**  $X' \in \mathbb{R}^{N \times K}$  (transformed data with reduced dimensions)
-

# Sample Covariance Matrix

- Eigen-vectors of symmetric matrices are orthogonal
- Covariance matrix is symmetric
  - Principal component are orthonormal
- We have


$$v_i^T v_j = 0, \quad \forall i \neq j$$

$$v_i^T v_i = 1, \quad \forall i$$

- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs**
- 4 Applications
- 5 Shortcomings
- 6 Conclusion
- 7 References

# How many PCs?

- For  $n$  original dimensions, sample covariance matrix is  $n * n$ , and has up to  $n$  eigenvectors. So  $n$  PCs
- Can ignore the components of lesser significance

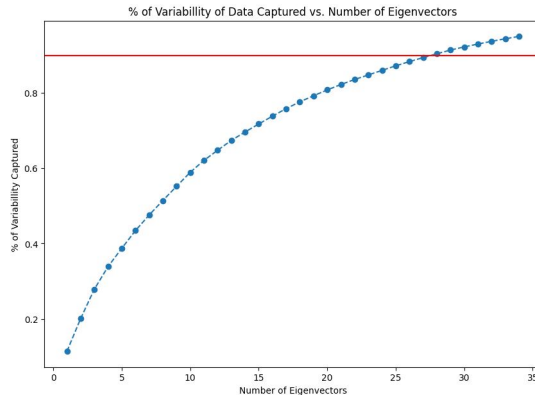


- You do lose some information, but if the eigenvalues are small, you don't lose much



# How many PCs?

$$\frac{\sum_{i=d-k+1}^d \lambda_i}{\sum_{i=1}^d \lambda_i}$$



- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs
- 4 Applications**
- 5 Shortcomings
- 6 Conclusion
- 7 References

# Image Compression

- Divide the original 372x492 image into patches
  - Each patch is an instance that contains 12x12 pixels on a grid
- Consider each as a 144-D vector



# Image Compression

- $144D \Rightarrow 60D$




# Image Compression

- $144D \Rightarrow 16D$



# Image Compression

- 16 most important eigenvectors






# Image Compression

- $144\text{D} \Rightarrow 3\text{D}$



# Image Compression


- 3 most important eigenvectors





# Image Compression


- L2 error and PCA dim



- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs
- 4 Applications
- 5 Shortcomings**
- 6 Conclusion
- 7 References


# Class Labels

- PCA doesn't know about class labels!



# Non-Linear

- PCA cannot capture Non-Linear structure!



- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs
- 4 Applications
- 5 Shortcomings
- 6 Conclusion
- 7 References

# Conclusion

- PCA
  - finds orthonormal basis for data
  - Sorts dimensions in order of “importance”
  - Discard low significance dimensions
- Applications
  - Get compact description
  - Remove noise
  - Improve classification (hopefully)
  - More efficient use of resources
  - Statistical
- Not magic
  - Doesn't know class labels
  - Can only capture linear variations
- One of many tricks to reduce dimensionality!

- 1 Introduction
- 2 Principal Component Analysis (PCA)
- 3 Choose PCs
- 4 Applications
- 5 Shortcomings
- 6 Conclusion
- 7 References

# References

- Advanced Introduction to Machine Learning, CMU-10715, Barnabás Póczos
- 10-701 Introduction to Machine Learning, CMU, Matt Gormley
- 10-301/10-601 Introduction to Machine Learning, CMU, Matt Gormley
- CE-477: Machine Learning - CS-828: Theory of Machine Learning, Sharif University of Technology, Fatemeh Seyyedsalehi