


Introduction to Neural Networks

ML Instruction Team, Fall 2022

CE Department
Sharif University of Technology



Biological Analogy




Figure: Anatomy of a biological neuron [1].

Activation Functions

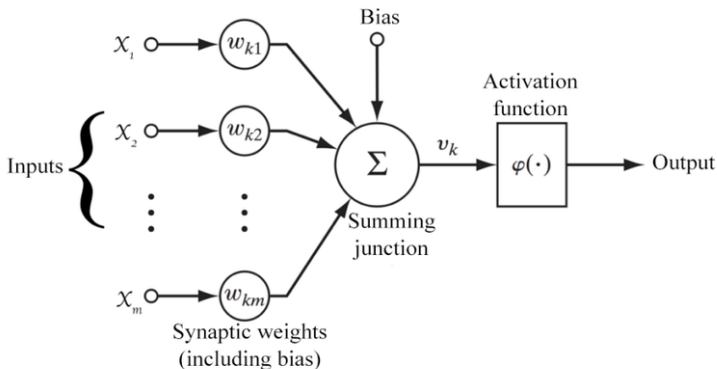


Figure: Neural network neuron [1].

Activation Functions

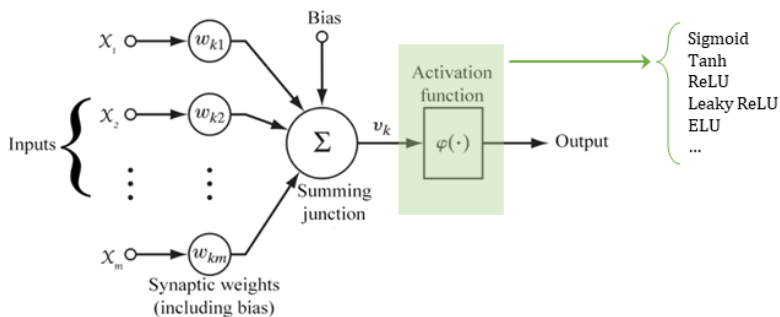



Figure: Activation function

Activation Functions


ReLU

$$f(x) = \begin{cases} x & x > 0 \\ 0 & x \leq 0 \end{cases}$$




ELU

$$f(x) = \begin{cases} x & x > 0 \\ \alpha(e^x - 1) & x \leq 0 \end{cases}$$



Leaky ReLU


$$f(x) = \begin{cases} x & x \geq 0 \\ 0.01x & x < 0 \end{cases}$$



Activation Functions


Tanh

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$




Sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



GELU

$$f(x) = \frac{1}{2}x \left(1 + \operatorname{erf}\left(\frac{x}{\sqrt{2}}\right) \right)$$

Softmax

$$f(x) = \frac{e^{x_i}}{\sum_{j=1}^J e^{x_j}} \quad i = 1, \dots, J$$

Gradient Descent




Figure: Gradient descent [3].

Gradient Descent

■ Let's define our problem:

- ▶ We have dataset $\mathcal{D} = \{x^i, y^{(i)}\}_{i=1}^n$.
- ▶ f is a single layer perceptron.
- ▶ Define $\hat{y}^{(i)} = f(x^{(i)})$.

■ We want to minimize following cost function:

$$\mathcal{J}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$

■ We are going to use gradient descent algorithm. \mathbf{w} will be updated as follows:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla_{\mathbf{w}} \mathcal{J}$$

Gradient Descent

■ Let's find $\nabla_{\mathbf{w}} \mathcal{J}$:

$$\begin{aligned}\frac{\partial J}{\partial w_j} &= \frac{\partial}{\partial w_j} \frac{1}{2} \sum_i (y^{(i)} - \hat{y}^{(i)})^2 \\&= \frac{1}{2} \sum_i \frac{\partial}{\partial w_j} (y^{(i)} - \hat{y}^{(i)})^2 \\&= \frac{1}{2} \sum_i 2(y^{(i)} - \hat{y}^{(i)}) \frac{\partial}{\partial w_j} (y^{(i)} - \hat{y}^{(i)}) \\&= \sum_i (y^{(i)} - \hat{y}^{(i)}) \frac{\partial}{\partial w_j} \left(y^{(i)} - \sum_j w_j x_j^{(i)} \right) \\&= \sum_i (y^{(i)} - \hat{y}^{(i)}) (-x_j^{(i)})\end{aligned}$$

$$\Delta w_j = -\eta \frac{\partial J}{\partial w_j} = -\eta \sum_i (y^{(i)} - \hat{y}^{(i)}) (-x_j^{(i)}) = \eta \sum_i (y^{(i)} - \hat{y}^{(i)}) x_j^{(i)}$$

$$\mathbf{w} := \mathbf{w} + \Delta \mathbf{w}$$

Limitations of SLP

- What are the limitations of SLP?
- Can we learn all functions with SLP?




Figure: The XOR function is not linear separable.

Limitations of SLP

- As we saw in the XOR case, nonlinear separable functions can not be learned by SPLs.




Figure: Examples of nonlinear separable functions.

- How to solve this?

Limitations of SLP

- What if we knew some feature space which our data is linear separable in?!




Figure: Data is linear separable after transformation.

Multi-Layer Perceptron

- So if we know some f_1, \dots, f_4 we can use SLP to solve problem.




Figure: Using feature space f to solve problem.

- How to learn this f_i s? Use SLP!

Multi-Layer Perceptron

- We can use inputs (x_i) to learn features (f_i)




Figure: Using inputs to learn features.

- What if f_i s are not sufficient? We can add more layers!

Multi-Layer Perceptron

- Adding more layers we will have a bigger network.




Figure: A 5 layer MLP.

Architecture of MLPs

■ Important questions:

- ▷ How many hidden layer should we have?
- ▷ In each hidden layer, how many neuron should we have?




Figure: How many layers and neurons should we have?

Architecture of MLPs

■ In practice:

- ▷ You have limited resources
- ▷ There is no universal rule to choose this hyperparameters
- ▷ Need to experiment for different number of layers and neurons in each layer




Figure: Experiment for different architecture and choose the best model.

Activation Function of Hidden Layers

- One can use any activation function for each hidden units
- Usually people use the same activation function for all neurons in one layer
- The important point is to use **nonlinear** activation functions




Figure: MLP with linear activation functions is equivalent to simple SLP.

XOR problem

- Now let's solve XOR problem with MLPs.
- We have two binary inputs, build an MLP to calculate their **XOR**.
- First let's build logical **AND** and **OR** functions.






Figure: We need to find weights, biases and activation function.


XOR problem



(a) $x_1 \wedge x_2$



(b) $x_1 \vee x_2$



(c) $x_1 \wedge \overline{x_2}$

XOR problem





Figure: MLP for XOR function.


MLP notation

- $a_i^{[l]}$: i -th neuron output in layer l
- $\mathbf{a}^{[l]}$: layer l output in vector form




MLP notation

- $b_i^{[l]}$: i -th neuron bias in layer l
- $\mathbf{b}^{[l]}$: layer l biases in vector form




MLP notation

- $W_{ij}^{[l]}$: weight of the edge between i -th neuron in layer $l - 1$ and j -th neuron in layer l




MLP notation

- $z_j^{[l]}$: j -th neuron input in layer l
- $z_j^{[l]} = b_j^{[l]} + \sum_{i=1}^n W_{ij}^{[l]} a_i^{[l-1]}$




MLP notation

- $\mathbf{z}^{[l]}$: input of layer l in vector form
- $\mathbf{z}^{[l]} = \mathbf{b}^{[l]} + (W^{[l]})^T \mathbf{a}^{[l-1]}$



MLP notation


- $\sigma_j^{[l]}$: j -th neuron activation function in layer l



MLP notation


- If we assume all neurons in one layer have the same activation function then:

$$\mathbf{a}^{[l]} = \sigma^{[l]} \left(\mathbf{b}^{[l]} + (W^{[l]})^T \mathbf{a}^{[l-1]} \right)$$



MLP notation


- So for a network with L layer, and \mathbf{x} as its input we will have:



$$\mathbf{o} = \mathbf{a}^{[L]} = \sigma^{[L]} \left(\mathbf{b}^{[L]} + (W^{[L]})^T \sigma^{[L-1]} \left(\dots \sigma^{[1]} \left(\mathbf{b}^{[1]} + (W^{[1]})^T \mathbf{x} \right) \dots \right) \right)$$

Learning MLPs

- Till here we have used networks with predefined weights and biases.
- How to learn these parameters?
- The idea is to use gradient descent



Learning MLPs

■ Let's define the learning problem more formal:

- ▷ $\{(x^{(i)}, y^{(i)})\}_{i=1}^n$: dataset
- ▷ f : network
- ▷ W : all weights and biases of the network ($W^{[l]}$ and $b^{[l]}$ for different l)
- ▷ L : loss function

■ We want to find W^* which minimizes following cost function:

$$\mathcal{J}(W) = \sum_{i=1}^n L\left(f(x^{(i)}; W), y^{(i)}\right)$$

■ We are going to use gradient descent, so we need to find $\nabla_W \mathcal{J}$.

Forward Propagation

- First of all we need to find loss value.
- It only requires to know the inputs of each neuron.




Figure:
$$a_j^{[l]} = \sum_{i=1}^m W_{ij}^{[l]} a_i^{[l-1]} + b_j^{[l]}$$

- So we can calculate these outputs layer by layer.

Forward Propagation

■ After forward pass we will know:

- ▷ Loss value
- ▷ Network output
- ▷ Middle values




Figure: Forward pass

Backward Propagation

- Now we need to calculate $\nabla_W \mathcal{J}$.
- First idea:
 - ▷ Use analytical approach.
 - ▷ Write down derivatives on paper.
 - ▷ Find the close form of $\nabla_W \mathcal{J}$ (if it is possible to do so).
 - ▷ Implement this gradient as a function to work with.
- ▷ Pros:
 - Fast
 - Exact
- ▷ Cons:
 - Need to rewrite calculation for different architectures

Backward Propagation

■ Second idea:

- ▷ Using **modular** approach.
- ▷ Computing the cost function consists of doing many operations.
- ▷ We can build a computation graph for this calculation.
- ▷ Each node will represent a single operation.





Figure: An example of computational graph, [Source](#).

Backward Propagation

- In this approach if we know how to calculate gradient for single node or module, then we can find gradient with respect to each variables.
- Let's say we have a module as follow:



- It gets x and y as its input and returns $z = f(x, y)$ as its output.
- How to calculate derivative of loss with respect to module inputs?


Backward Propagation

■ We know:

- ▷ Module output for x_0 and y_0 , let's call it z_0 .
- ▷ Gradient of loss with respect to module output at z_0 , $\left(\frac{\partial L}{\partial z}\right)$.

■ We want:

- ▷ Gradient of loss with respect to module inputs at x_0 and y_0 , $\left(\frac{\partial L}{\partial x}, \frac{\partial L}{\partial y}\right)$.




Backward Propagation

- We can use chain rule to do so.


Chain rule:

$$\left. \begin{array}{l} z = f(x, y) \\ L = L(z) \end{array} \right\} \Rightarrow \frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \times \frac{\partial z}{\partial x}$$



Backward Propagation

■ Backpropagation for single module:



Backward Propagation: Example

- Let's solve a simple example using backpropagation.
- We have $f(x, y, z) = \frac{x^2 y}{z}$.
- Find $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$ and $\frac{\partial f}{\partial z}$ at $x = 3$, $y = 4$ and $z = 2$.




Figure: Computational graph of f .

Backward Propagation: Example


- First let's find gradient analytically.
- We have:

$$\begin{cases} v = x^2 \\ u = vy \\ f = \frac{u}{z} \end{cases} \quad \begin{cases} \frac{\partial f}{\partial z} = -\frac{u}{z^2} \\ \frac{\partial f}{\partial y} = \frac{\partial u}{\partial y} \times \frac{\partial f}{\partial u} = v \times \frac{1}{z} = \frac{v}{z} \\ \frac{\partial f}{\partial x} = \frac{\partial v}{\partial x} \times \frac{\partial u}{\partial v} \times \frac{\partial f}{\partial u} = 2x \times y \times \frac{1}{z} = \frac{2xy}{z} \end{cases}$$

$$(x = 3, y = 4, z = 2) \implies \begin{cases} v = 9 \\ u = 36 \\ f = 18 \end{cases} \implies \begin{cases} \frac{\partial f}{\partial z} = -9 \\ \frac{\partial f}{\partial y} = 4.5 \\ \frac{\partial f}{\partial x} = 12 \end{cases}$$

Backward Propagation: Example


- Now let's use backpropagation.
- First we do forward propagation.



- Second we will do backpropagation for each module.


Backward Propagation: Example

Backpropagation for / module:




Backward Propagation: Example

Backpropagation for / module:




Backward Propagation: Example

Backpropagation for \times module:




Backward Propagation: Example

Backpropagation for \times module:




Backward Propagation: Example

Backpropagation for $_2$ module:



Backward Propagation: Example

Backpropagation for $_2^2$ module:



Results are the same as analytical results.

Backward Propagation

- So after backward propagation we will have:
 - ▷ Gradient of loss with respect to each parameter.
 - ▷ We can apply gradient descent to update parameters.




Figure: Backward pass

Thank You!

Any Question?

References



M. D. Ergün Akgün, “Biological and neural network neuron,” 2018.

https://www.researchgate.net/publication/326417061_Modeling_Course_Achievements_of_Elementary_Education_Teacher_Candidates_with_Artificial_Neural_Networks.



L. Nalborczyk, “A gentle introduction to deep learning in r using keras,” 2021.

https://www.barelyssignificant.com/slides/vendredi_quanti_2021/vendredi_quantis#1.



“Gradient descent.”

<https://subscription.packtpub.com/book/big-data-&-business-intelligence/9781788397872/1/ch01lv1sec22/gradient-descent>.



F.-F. L. . J. J. . S. Yeung, “Training neural networks,” 2018.

http://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture07.pdf.



I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*.

MIT Press, 2016.

<http://www.deeplearningbook.org>.



K. Katanforoosh and D. Kunin, “Initializing neural networks,” 2019.

<https://www.deeplearning.ai/ai-notes/initialization/>.