



# Generative Adversarial Networks (GANs)

ML Instruction Team, Fall 2022



CE Department  
Sharif University of Technology

## Module 23.1: Generative Adversarial Networks - The intuition



- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution





- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution
- For example, in RBMs we learn  $P(X, H)$ , in VAEs we learn  $P(z|X)$  and  $P(X|z)$  whereas in AR models we learn  $P(X)$




- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution
- For example, in RBMs we learn  $P(X, H)$ , in VAEs we learn  $P(z|X)$  and  $P(X|z)$  whereas in AR models we learn  $P(X)$
- What if we are only interested in sampling from the distribution and don't really care about explicit density function  $P(X)$ ?



- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution
- For example, in RBMs we learn  $P(X, H)$ , in VAEs we learn  $P(z|X)$  and  $P(X|z)$  whereas in AR models we learn  $P(X)$
- What if we are only interested in sampling from the distribution and don't really care about explicit density function  $P(X)$ ?
- What does this mean?




- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution
- For example, in RBMs we learn  $P(X, H)$ , in VAEs we learn  $P(z|X)$  and  $P(X|z)$  whereas in AR models we learn  $P(X)$
- What if we are only interested in sampling from the distribution and don't really care about explicit density function  $P(X)$ ?
- What does this mean? Let us see




- As usual we are given some training data (say, MNIST images) which obviously comes from some underlying distribution







- As usual we are given some training data (say, MNIST images) which obviously comes from some underlying distribution
- Our goal is to generate more images from this distribution (*i.e.*, create images which look similar to the images from the training data)



- As usual we are given some training data (say, MNIST images) which obviously comes from some underlying distribution
- Our goal is to generate more images from this distribution (*i.e.*, create images which look similar to the images from the training data)
- In other words, we want to sample from a complex high dimensional distribution which is intractable (recall RBMs, VAEs and AR models deal with this intractability in their own way)



- GANs take a different approach to this problem where the idea is to sample from a simple tractable distribution (say,  $z \sim N(0, I)$ ) and then learn a complex transformation from this to the training distribution



- GANs take a different approach to this problem where the idea is to sample from a simple tractable distribution (say,  $z \sim N(0, I)$ ) and then learn a complex transformation from this to the training distribution
- In other words, we will take a  $z \sim N(0, I)$ , learn to make a series of complex transformations on it so that the output looks as if it came from our training distribution

- What can we use for such a complex transformation?

- What can we use for such a complex transformation? A Neural Network


- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network?

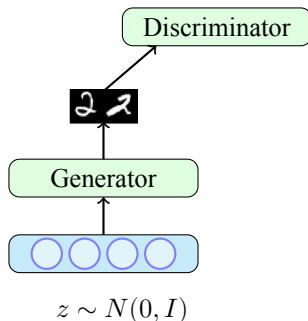
- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network?  
Using a two player game



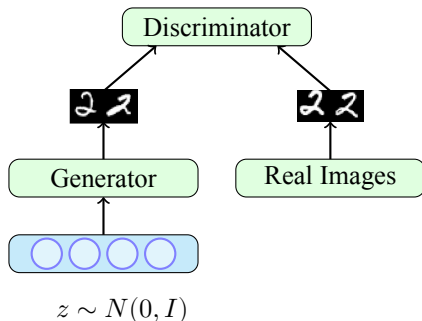
- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network?  
Using a two player game
- There are two players in the game:

- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a **generator**






- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and a **discriminator**

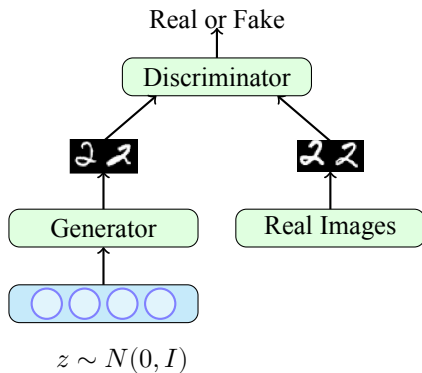



- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and a discriminator
- The job of the generator is to produce images which look so natural that the discriminator thinks that the images came from the real data distribution




- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and a discriminator
- The job of the generator is to produce images which look so natural that the discriminator thinks that the images came from the real data distribution
- The job of the discriminator is to get better and better at distinguishing between true images and generated (fake) images

- So let's look at the full picture






- So let's look at the full picture
- Let  $G_\phi$  be the generator and  $D_\theta$  be the discriminator ( $\phi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)



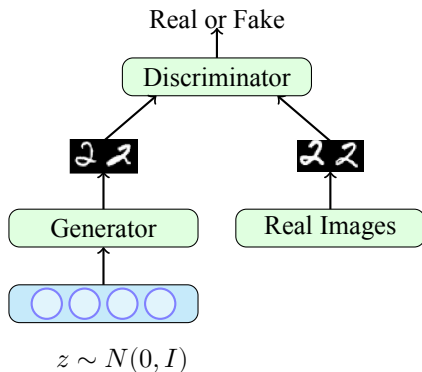
- So let's look at the full picture
- Let  $G_\phi$  be the generator and  $D_\theta$  be the discriminator ( $\phi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)
- We have a neural network based generator which takes as input a noise vector  $z \sim N(0, I)$  and produces  $G_\phi(z) = X$







- So let's look at the full picture
- Let  $G_\phi$  be the generator and  $D_\theta$  be the discriminator ( $\phi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)
- We have a neural network based generator which takes as input a noise vector  $z \sim N(0, I)$  and produces  $G_\phi(z) = X$
- We have a neural network based discriminator which could take as input a real  $X$  or a generated  $X = G_\phi(z)$  and classify the input as real/fake

- What should be the objective function of the overall network?







- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first



- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it





- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it
- This score will be between 0 and 1 and will tell us the probability of the image being real or fake



- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it
- This score will be between 0 and 1 and will tell us the probability of the image being real or fake
- For a given  $z$ , the generator would want to maximize  $\log D_\theta(G_\phi(z))$  (log likelihood) or minimize  $\log(1 - D_\theta(G_\phi(z)))$

- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,






- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,
- For example, if  $z$  was discrete and drawn from a uniform distribution (*i.e.*,  $p(z) = \frac{1}{N} \forall z$ ) then the generator's objective function would be

$$\min_{\phi} \sum_{i=1}^N \frac{1}{N} \log(1 - D_{\theta}(G_{\phi}(z)))$$





- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,
- For example, if  $z$  was discrete and drawn from a uniform distribution (*i.e.*,  $p(z) = \frac{1}{N} \forall z$ ) then the generator's objective function would be


$$\min_{\phi} \sum_{i=1}^N \frac{1}{N} \log(1 - D_{\theta}(G_{\phi}(z)))$$


- However, in our case,  $z$  is continuous and not uniform ( $z \sim N(0, I)$ ) so the equivalent objective function would be

$$\min_{\phi} \int p(z) \log(1 - D_{\theta}(G_{\phi}(z)))$$


$$\min_{\phi} E_{z \sim p(z)} [\log(1 - D_{\theta}(G_{\phi}(z)))]$$

- Now let's look at the discriminator






- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images




- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images
- And it should do this for all possible real images and all possible fake images




- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images
- And it should do this for all possible real images and all possible fake images
- In other words, it should try to maximize the following objective function

$$\max_{\theta} E_{x \sim p_{data}} [\log D_{\theta}(x)] + E_{z \sim p(z)} [\log(1 - D_{\theta}(G(z)))]$$



- If we put the objectives of the generator and discriminator together we get a minimax game


$$\min_{\phi} \max_{\theta} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z))) \right]$$



- If we put the objectives of the generator and discriminator together we get a minimax game

$$\min_{\phi} \max_{\theta} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z))) \right]$$

- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )




- If we put the objectives of the generator and discriminator together we get a minimax game

$$\min_{\phi} \max_{\theta} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z))) \right]$$

- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )
- The second term in the objective is w.r.t. the parameters of the generator ( $\phi$ ) as well as the discriminator ( $\theta$ )



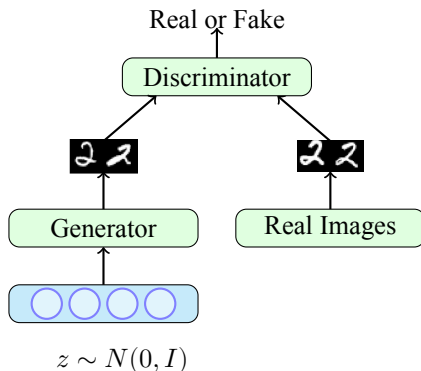



- If we put the objectives of the generator and discriminator together we get a minimax game

$$\min_{\phi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )
- The second term in the objective is w.r.t. the parameters of the generator ( $\phi$ ) as well as the discriminator ( $\theta$ )
- The discriminator wants to maximize the second term whereas the generator wants to minimize it (hence it is a two-player game)


- So the overall training proceeds by alternating between these two step





- So the overall training proceeds by alternating between these two step
- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(x))]$$




- So the overall training proceeds by alternating between these two step

- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- **Step 2:** Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$




- So the overall training proceeds by alternating between these two step
- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(z))]$$

- ### Step 2: Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$

- In practice, the above generator objective does not work well and we use a slightly modified objective



- So the overall training proceeds by alternating between these two step

- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- **Step 2:** Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$

- In practice, the above generator objective does not work well and we use a slightly modified objective
- Let us see why

**Thank You!**

**Any Question?**